

```
In [1]: import pandas as pd
```

```
In [2]: df = pd.read_csv("emails.csv")
```

```
In [3]: df.shape
```

```
Out[3]: (5172, 3002)
```

```
In [4]: df.head()
```

```
Out[4]:
```

	Email No.	the	to	ect	and	for	of	a	you	hou	...	connevey	jay	valued	lay	infrastructu
0	Email 1	0	0	1	0	0	0	2	0	0	...	0	0	0	0	
1	Email 2	8	13	24	6	6	2	102	1	27	...	0	0	0	0	
2	Email 3	0	0	1	0	0	0	8	0	0	...	0	0	0	0	
3	Email 4	0	5	22	0	5	1	51	2	10	...	0	0	0	0	
4	Email 5	7	6	17	1	5	2	57	0	9	...	0	0	0	0	

5 rows × 3002 columns



```
In [5]: x = df.drop(['Email No.', 'Prediction'], axis = 1)
        y = df['Prediction']
```

```
In [6]: x.shape
```

```
Out[6]: (5172, 3000)
```

```
In [7]: x.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5172 entries, 0 to 5171
Columns: 3000 entries, the to dry
dtypes: int64(3000)
memory usage: 118.4 MB
```

```
In [8]: x.dtypes
```


```
Out[8]: the          int64
to            int64
ect           int64
and           int64
for           int64
...
infrastructure int64
military       int64
allowing       int64
ff             int64
dry            int64
Length: 3000, dtype: object
```

```
In [9]: set(x.dtypes)
```

```
Out[9]: {dtype('int64')}
```

```
In [11]: import seaborn as sns
sns.countplot(x=y)
```

```
Out[11]: <Axes: xlabel='Prediction', ylabel='count'>
```



```
In [12]: y.value_counts()
```

```
Out[12]: 0    3672
1     1500
Name: Prediction, dtype: int64
```

```
In [15]: from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler()
x_scaled = scaler.fit_transform(x)
```

```
In [16]: x_scaled
```

```
Out[16]: array([[0.          , 0.          , 0.          , ..., 0.          , 0.          ,
                0.          ],
               [0.03809524, 0.09848485, 0.06705539, ..., 0.          , 0.00877193,
                0.          ],
               [0.          , 0.          , 0.          , ..., 0.          , 0.          ,
                0.          ],
               ...,
               [0.          , 0.          , 0.          , ..., 0.          , 0.          ,
                0.          ],
               [0.00952381, 0.0530303 , 0.          , ..., 0.          , 0.00877193,
                0.          ],
               [0.1047619 , 0.18181818, 0.01166181, ..., 0.          , 0.          ,
                0.          ]])
```

```
In [17]: from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(
x_scaled,y, random_state=0, test_size = 0.25)
```

```
In [18]: x_scaled.shape
```

```
Out[18]: (5172, 3000)
```

```
In [19]: x_train.shape
```

```
Out[19]: (3879, 3000)
```

```
In [21]: x_test.shape
```

```
Out[21]: (1293, 3000)
```

```
In [22]: from sklearn.neighbors import KNeighborsClassifier
```

```
In [23]: knn = KNeighborsClassifier(n_neighbors=5)
```

```
In [24]: knn.fit(x_train, y_train)
```

```
Out[24]: 

▼ KNeighborsClassifier



KNeighborsClassifier()
```

```
In [25]: y_pred = knn.predict(x_test)
```

```
In [27]: from sklearn.metrics import ConfusionMatrixDisplay, accuracy_score
from sklearn.metrics import classification_report
```

```
In [30]: ConfusionMatrixDisplay.from_predictions (y_test, y_pred)
```

```
Out[30]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x176d9d3d450>
```



```
In [32]: y_test.value_counts()
```

```
Out[32]: 0    929
         1    364
         Name: Prediction, dtype: int64
```

```
In [33]: accuracy_score(y_test, y_pred)
```

```
Out[33]: 0.871616395978345
```

```
In [34]: print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.98	0.84	0.90	929
1	0.70	0.95	0.81	364
accuracy			0.87	1293
macro avg	0.84	0.89	0.85	1293
weighted avg	0.90	0.87	0.88	1293

```
In [39]: import numpy as np
import matplotlib.pyplot as plt
```

```
In [41]: error = []
for k in range(1,41):
    knn = KNeighborsClassifier(n_neighbors=k)
    knn.fit(x_train, y_train)
    pred = knn.predict(x_test)
    error.append(np.mean(pred != y_test))
```

```
In [42]: error
```

```
Out[42]: [0.10827532869296211,
0.10982211910286156,
0.12296983758700696,
0.11523588553750967,
0.12838360402165508,
0.1214230471771075,
0.15158546017014696,
0.14849187935034802,
0.17246713070378963,
0.16705336426914152,
0.1871616395978345,
0.18329466357308585,
0.21500386697602475,
0.21345707656612528,
0.22815158546017014,
0.2266047950502707,
0.23588553750966745,
0.23356535189481825,
0.2459396751740139,
0.24361948955916474,
0.2559938128383604,
0.2552204176334107,
0.2699149265274555,
0.2691415313225058,
0.2822892498066512,
0.28306264501160094,
0.2954369682907966,
0.2923433874709977,
0.3039443155452436,
0.300077339520495,
0.30549110595514306,
0.30549110595514306,
0.31245166279969067,
0.31245166279969067,
0.3194122196442382,
0.317092034029389,
0.32637277648878577,
0.32559938128383603,
0.33410672853828305,
0.3325599381283836]
```

```
In [43]: knn = KNeighborsClassifier(n_neighbors=1)
```

```
In [44]: knn.fit(x_train,y_train)
```

```
Out[44]: 

▼



KNeighborsClassifier



KNeighborsClassifier(n_neighbors=1)


```

```
In [45]: y_pred = knn.predict(x_test)
```

```
In [47]: accuracy_score(y_test, y_pred)
```

```
Out[47]: 0.8917246713070379
```

```
In [48]: from sklearn.svm import SVC
```

```
In [49]: svm = SVC(kernel= 'linear')
```

```
In [50]: svm.fit(x_train, y_train)
```

```
Out[50]: 

▼



SVC



SVC(kernel='linear')


```

```
In [51]: y_pred = svm.predict(x_test)
```

```
In [53]: accuracy_score(y_test, y_pred)
```

```
Out[53]: 0.9767981438515081
```

```
In [ ]:
```